

# Lab 1 - Descriptive Statistics

January 24, 2020

## 1 Packages and Data Structures

The Python core distribution contains only the functions and methods essential for a general programming language. The power of Python lies in its modules and packages. These are collections of functions that can be imported into and used in your Python workspace. Throughout this module we will use many packages, the most frequent being NumPy, SciPy and Pandas. \* **NumPy** allows for the creation of vectors, matrices and higher-dimensional arrays. It allows contains functions for performing operations on these objects, such as matrix multiplication and solving linear systems. \* **SciPy** builds on the functionalities of NumPy, with modules for optimisation, statistics, numerical integration, differential equations and linear algebra. There is a lot of overlap between NumPy and SciPy. \* **Pandas** is designed for elegant data manipulation and analysis. It allows for data to be structured as tables and contains functions for performing operations on these tables.

Packages need to be imported at the beginning of each script. Typically, we import packages with an abbreviation so that we don't need to type out the full package name everytime we wish to use a function from the package.

```
[1]: import numpy as np
import scipy as sp
import pandas as pd
```

NumPy *arrays* can be created from basic Python objects, such as lists, tuples and dicts, using the `np.array()` function. If we perform a standard operation, such as addition or multiplication, to the array it will be applied to each element of the array

```
[2]: a = [2,5,7,1,6,3]
vec = np.array(a)
print(vec+1)
print(vec*2)
print(vec[2:4])
```

```
[3 6 8 2 7 4]
[ 4 10 14  2 12  6]
[7 1]
```

As mentioned above Pandas provides data structures for elegant data manipulation and analysis. These data structures are known as *DataFrames*, and are similar to spreadsheets. The standard way to create a DataFrame is from a dict or by loading in a data file as a DataFrame. It is convenient

to import the DataFrame function from Pandas to save using the `pd.` prefix every time we want to use it.

```
[3]: from pandas import DataFrame
dict1 = {'Names': ['Aoife', 'Brian', 'Catherine', 'Daniel', 'Eamonn', 'Fiona'],
        'Age': [23, 45, 17, 64, 57, 32], 'Height': [1.7, 1.9, 1.55, 1.8, 1.75, 1.65],
        'Weight': [82, 88, 55, 101, 75, 67]}
df = DataFrame(dict1)
print(df)
```

	Names	Age	Height	Weight
0	Aoife	23	1.70	82
1	Brian	45	1.90	88
2	Catherine	17	1.55	55
3	Daniel	64	1.80	101
4	Eamonn	57	1.75	75
5	Fiona	32	1.65	67

Columns are referenced by their heading and can be accessed using either `df.ColumnName` or `df['ColumnName']`, e.g. `df.Names` or `df['Names']` for the first column. You can also access elements in the DataFrame using their position with `df.iloc[a,b]`, where `a` and `b` are the row and column numbers, respectively, that you wish to access. Remember that indexing in Python starts at 0, so `df.iloc[4,2]` would return the entry in the 5th row and 3rd column, i.e. Eamonn's height 1.75.

```
[4]: print(df.iloc[3:5,:2])
```

	Names	Age
3	Daniel	64
4	Eamonn	57

## 2 Importing Data

There are many ways to import data in Python. We will predominantly use the Pandas `pd.read_csv()` function to load data files into Python as a DataFrame. Download the dataset `marks.csv` from Brightspace and save it into your current working directory. To useful functions for inspecting DataFrames are `.head()` and `.tail()`. They return the first 5 entries and last 5 entries, respectively.

```
[5]: data = pd.read_csv('marks.csv')
print(data.head())           # printing all 365 entries would be ridiculous
```

	StudentID	Coursework	Project	Exam
0	34258	86.4	68	90.67
1	566811	77.4	55	90.67
2	6359256	15.5	66	40.00
3	6361307	83.3	85	92.00
4	6390081	85.2	90	93.33

### 3 Numerical summaries

Pandas DataFrames have lots of associated methods and function. If you type `data.` into your Python shell and hit the tab key you will get a list of possible completions, these are all of the methods that can applied to `data`, as well as the objects associated with it. `data.describe()` will produce a table with summary statistics for each of the columns in `data`.

```
[6]: print(data.describe())
```

	Coursework	Project	Exam
count	51.000000	51.000000	51.000000
mean	76.807843	73.333333	70.535882
std	20.465921	13.387556	19.616650
min	15.500000	35.000000	24.000000
25%	65.550000	65.000000	53.330000
50%	85.200000	75.000000	74.670000
75%	91.300000	85.000000	86.670000
max	98.600000	95.000000	96.000000

We can compute these statistics individually with functions such as `mean()`, `std()`, `min()` etc.

```
[7]: print(data.mean())
std_exam = data['Exam'].std()
print('The standard deviation of the exam score was', std_exam, '%.')
max_proj = data['Project'].max()
max_proj_ID = data.loc[data['Project'].idxmax(), 'StudentID']
print('The best project score was achieve by student', max_proj_ID, 'with a
↪score of'
      , max_proj, '%.')
```

```
Coursework    76.807843
Project       73.333333
Exam          70.535882
dtype: float64
The standard deviation of the exam score was 19.616650088786376 %.
The best project score was achieve by student 6900976 with a score of 95 %.
```

### 4 Graphical summaries

The best plotting package in Python is Matplotlib, which contains functions for all of the plots we will consider in this lab session. For more examples and information on Matplotlib see the [documentation page](#).

```
[8]: import matplotlib.pyplot as plt

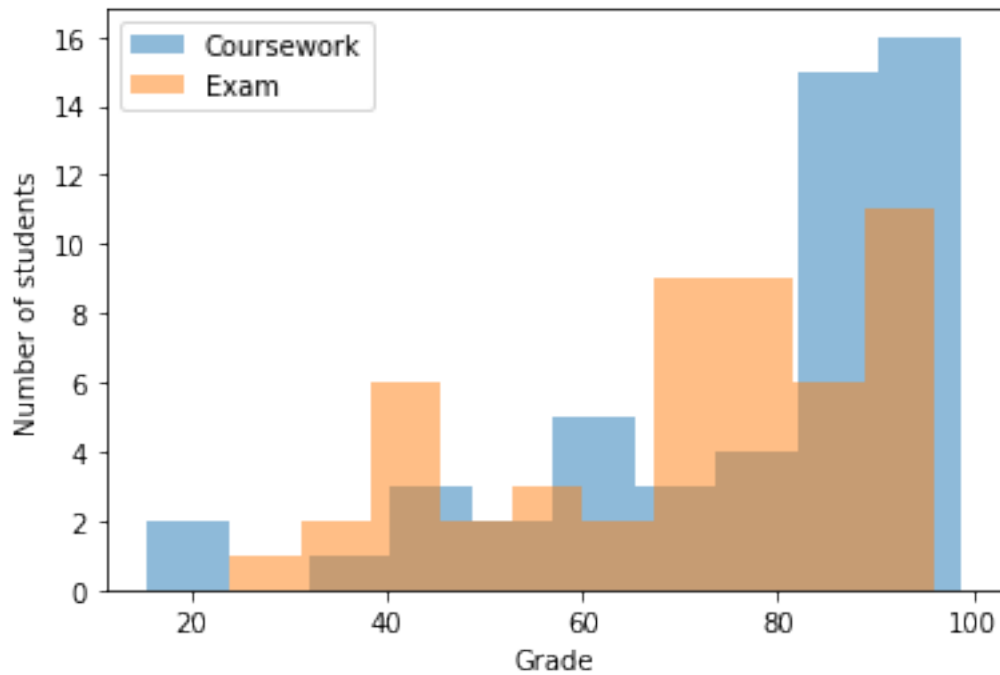
plt.figure()
plt.scatter(data.Coursework, data.Exam)
```

```
plt.xlabel('Coursework grade')
plt.ylabel('Exam grade')
plt.axis([0,100,0,100])
```

[8]: [0, 100, 0, 100]


```
[9]: plt.figure()
plt.hist(data.Coursework,bins=10,alpha=0.5,label='Coursework')
plt.hist(data.Exam,bins=10,alpha=0.5,label='Exam')
plt.xlabel('Grade')
plt.ylabel('Number of students')
plt.legend()
```

[9]: <matplotlib.legend.Legend at 0x1170d2790>



```
[10]: plt.figure()
data.boxplot()
plt.ylabel('Grade')
plt.ylim(0,100)
```

[10]: (0, 100)



## 5 Exercises

1. Download the `weather2019.csv` dataset from Brightspace and save it into your current working directory.
2. Start a **new** Jupyter notebook, import the necessary packages and load in the `weather2019.csv` dataset.
3. Go to the *Lab 1* Quiz on Brightspace to find your questions for this week. Note that your questions will be different to other students.
4. You will need to write Python code to answer the questions. This code should be submitted to the *Lab 1* Assignment object on Brightspace in addition to completing the Quiz.

**Note:** You should not submit your Quiz attempt unless you are completely finished. Starting a new attempt will result in a new set of questions. You can save your current attempt and come back to it later if you need to.

The deadline for submitting your Quiz attempt and Python file is **Monday February 3rd at noon**.

## 6 Additional non-assessed exercises

1. Download the `MichelsonMorley.csv` dataset from Brightspace and save it into your current working directory. This is the dataset from Michelson and Morley's experiments on the speed of light.
2. Reproduce the histogram on slide 19 of the lecture notes on Descriptive Statistics.

3. Interpret your plot. What do you learn about the data from it?
4. Describe what happens when you change the number of bins.
5. What is the mean measurement for the speed of light?
6. What is the standard deviation?
7. Explore the options associated with the histogram command (`plt.hist?`). Try to change the colour, transparency, labels, etc of your plot.